

# Lecture 11. Exploration strategies. RL outside games.

Nikolay Karpachev  
15.04.2024

# Outline

- Exploration / exploitation tradeoff
- Exploration strategies
  - Heuristic-based
  - “Optimism in the face of uncertainty”
  - Probability matching
- Structured prediction in NLP -> RL framing
- Optimizing for undifferentiable with policy gradient
- SCST

# Exploration strategies

# Exploration vs. Exploitation in RL

- Online decision-making involves a fundamental choice
  - **Exploitation**: Make the best decision given the current information
  - **Exploration**: Gather more information
- The best long-term strategy may involve short-term sacrifices
- Agent should gather enough relevant information to make reasonable decisions

# Exploration vs. Exploitation: examples

- Restaurant selection
  - **Exploitation**: Go to your favourite restaurant
  - **Exploration**: Try new restaurant
- Online banner advertisements
  - **Exploitation**: Show the most successful advert
  - **Exploration**: Show a different advert
- Game playing
  - **Exploitation**: Play the move you believe is the best
  - **Exploration**: Play a different move

# Multi-armed bandit

- What is a bandit?

# Multi-armed bandit



# Multi-armed bandit

- A single state
- Set of possible actions (decide which slot machine to play)
- Each machine has an unknown probability of success
- Goal: maximize the total number of successful games



# Regret

$$Q(a) = \mathbb{E}[r|a]$$

$$V^* = Q(a^*) = \max_{a \in \mathcal{A}} Q(a)$$

- Regret (Total Regret): opportunity loss for one step (all steps)

$$l_t = \mathbb{E}[V^* - Q(a_t)]$$

$$L_t = \mathbb{E}\left[\sum_{\tau=1}^t V^* - Q(a_\tau)\right]$$

**We want to minimize the total regret**

# Exploration strategies so far

- Eps-greedy
  - With  $p = \epsilon$ , take random action. Optimal otherwise
- Boltzman (aka softmax)
  - Pick actions proportionately to scaled Q-values

$$P(a) = \text{softmax}\left(\frac{Q(s, a)}{\text{std}}\right)$$

- Decaying eps-greedy
  - Same as eps-greedy; start with high eps, decrease it during training

# Greedy algorithm

- Always selects actions with highest values
- What is the total regret?

# Greedy algorithm

- Always selects actions with highest values
  - What is the total regret?
- 
- Greedy can lock to a suboptimal action forever
  - Hence, **linear total regret**

# Epsilon-greedy algorithm

- Explores forever
- Selects suboptimal actions with fixed probability over and over again
- Linear total regret

# Epsilon-greedy with decay

- Has a decay schedule for  $\epsilon$
- With properly selected schedule, has a **logarithmic total regret**
- But to design a proper schedule can be tricky

# Optimism in the face of uncertainty


- How do humans explore?
- For example, which of the following questions would you like to investigate?
  - Whether humans can fly by pulling their hair up
  - Whether the new cafe next to the office serves good breakfast

# Optimism in the face of uncertainty

- How do humans explore?
- For example, which of the following questions would you like to investigate:
  - Whether humans can fly by pulling their hair up
  - Whether the new cafe next to the office serves good breakfast
- We want to try actions if we believe there's a chance they are good

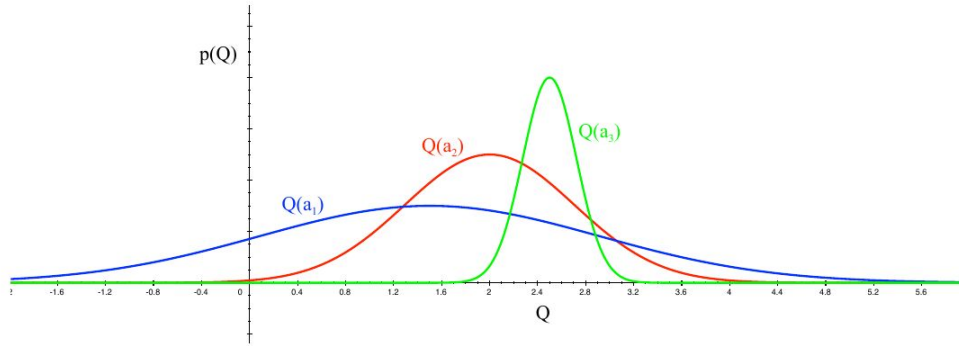


# Optimism in the face of uncertainty




- Which action should we pick?

# Optimism in the face of uncertainty



- Which action should we pick?
- The more uncertain we are about an action value
- The more important it is to try that action
- It could turn out to be the best action

# Optimism in the face of uncertainty



- After picking blue action
- We are less uncertain about the value
- And more likely to pick another action

# Upper confidence bounds

- We want to select
  - Uncertain outcomes
  - With greater expected value

# Upper confidence bounds

- We want to select
  - Uncertain outcomes
  - With greater expected value
- Let's compute 95% upper confidence bound for each action
- Take action with the highest upper confidence bound

# Upper confidence bounds

## Theorem (Hoeffding's inequality):

Given a sample of a random variable bounded in  $[0, 1]$ ,

$$\mathbb{P} [\mathbb{E} [X] > \overline{X}_t + u] \leq e^{-2tu^2}$$

# Upper confidence bounds

- We can apply Hoeffding's inequality to the case of bandits:

$$\mathbb{P} \left[ Q(a) > \hat{Q}_t(a) + U_t(a) \right] \leq e^{-2N_t(a)U_t(a)^2}$$

$$e^{-2N_t(a)U_t(a)^2} = p$$

$$U_t(a) = \sqrt{\frac{-\log p}{2N_t(a)}}$$

# Upper confidence bounds

- With fixed  $p$  (e.g. 95% UCB)

$$U_t(a) = \sqrt{\frac{-\log p}{2N_t(a)}}$$

- Possible extension: reduce  $p$  during training (ucb converges to 1 in the limit; this guarantees optimal solution)

$$p = t^{-4}$$

$$U_t(a) = \sqrt{\frac{2 \log t}{N_t(a)}}$$



# UCB-1 algorithm

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q(a) + \sqrt{\frac{2 \log t}{N_t(a)}}$$

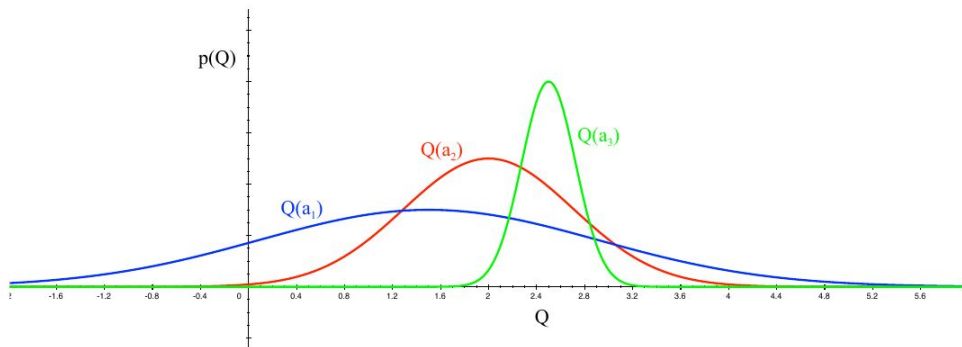
- Achieves **logarithmic total regret**

# Bayesian UCB

- Assign prior distribution  $P(Q(s,a))$
- Learn posterior  $P(Q(s,a)|\text{data})$
- Take  $q$ -th percentile of  $P(Q(s,a))$  and select the best action

# Probability matching

- Select action  $a$  according to the probability that  $a$  is the optimal action



$$\pi(a \mid h_t) = \mathbb{P} [Q(a) > Q(a'), \forall a' \neq a \mid h_t]$$

# Thompson sampling

- Compute posterior distribution for each  $Q(s,a)$
- Sample from each action's posterior
- Select action with max value on sample
- **Thompson sampling will select action proportionately to the probability that this action is optimal**

RL outside games

# NLP Training / Inference Regimes

- Supervised seq2seq learning:

$$P(y_{t+1}|x, y_{0:t}), \quad y_{0:t} \sim \text{reference}$$

- Inference

$$P(y_{t+1}|x, \hat{y}_{0:t}), \quad \hat{y}_{0:t} \sim ???$$

# NLP Training / Inference Regimes

- Supervised seq2seq learning:

$$P(y_{t+1}|x, y_{0:t}), \quad y_{0:t} \sim \text{reference}$$

- Inference

$$P(y_{t+1}|x, \hat{y}_{0:t}), \quad \hat{y}_{0:t} \sim ???$$

**If model ever makes something that isn't in data,  
It gets volatile from next time-step!**

# NLP Training / Inference Regimes

- Supervised seq2seq learning:

$$P(y_{t+1}|x, y_{0:t}), \quad y_{0:t} \sim \text{reference}$$

- Inference

$$P(y_{t+1}|x, \hat{y}_{0:t}), \quad \hat{y}_{0:t} \sim ???$$

“Exposure bias”



**If model ever makes something that isn't in data,  
It gets volatile from next time-step!**



# Issue 1: Exposure bias

- Need good training data
  - Abundant
  - Few biases and noise
- Need a perfect network to extrapolate training set

# Issue 1: Exposure bias

- Need good training data
  - Abundant
  - Few biases and noise
- Need a perfect network to extrapolate training set

Spoiler: Most of the time we don't. Too bad.



## Issue 2: Crossentropy loss

Reality: sometimes there is more than one correct translation (image caption, summary, whatever)

**Source:** 在找给家里人的礼物.

**Versions:**

i 'm searching for some gifts for my family.

i want to find something for my family as presents.

i 'm about to buy some presents for my family.

i 'd like to buy my family something as a gift.

i 'm looking for a present for my family.

...

## Issue 2: Crossentropy loss

Reality: sometimes there is more than one correct translation (image caption, summary, whatever)

**Source:** 在找给家里人的礼物.

**Versions:**

(version 1)  
(version 2)  
(version 3)  
(all rubbish)

not in data

Model 1  
 $p(y|x)$

1e-2  
2e-2  
1e-2  
0.96

Model 2  
 $p(y|x)$

0.99  
1e-100  
1e-100  
0.01

**Question:**  
which model  
has better  
Mean log  
 $p(y|x)$  ?

**This one. While it predicts 96% rubbish**

# Issue 3: Training data

Two kinds of datasets:

Big enough, but suboptimal  $R(x,y)$

- **Large raw data**

- twitter, open subtitles, books, bulk logs
- $10^6$ - $8$  samples, <http://opus.nlpl.eu/OpenSubtitles.php>

- **Small clean data**

- moderated logs, assessor-written conversations
- $10^2$ ~ $4$  samples

Near-optimal  $R(x,y)$ , but too small


# Can we do better?

Key idea: frame structured prediction problem as RL task

Pros:

- Can optimize any objective function **directly** (not just probability of correct stuff)
  - User feedback (discrete scores)
  - External reward model
  - Practically, any black-box stuff
- Inference data is the new training data!

# Seq2Seq as POMDP



*Hidden state  $\mathbf{s}$*  = translation/conversation state

Initial state  $\mathbf{s}$  = encoder output

Observation  $\mathbf{o}$  = previous words

Action  $\mathbf{a}$  = write next word

Reward  $\mathbf{r}$  = domain-specific reward (e.g. BLEU)

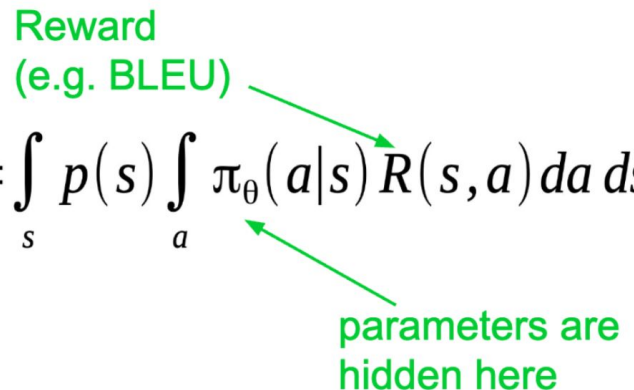
# Training with Policy Gradient

Our objective:

$$J = E_{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}} R(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds$$

Reward  
(e.g. BLEU)

parameters are hidden here



We can approximate the expectation with mean:

$$J \approx \frac{1}{N} \sum_{i=0}^N R(s, a)$$



# Training with Policy Gradient

Our objective:

$$J = \mathop{E}_{\substack{s \sim p(s) \\ a \sim \pi_\theta(s|a)}} R(s, a) = \int_s p(s) \int_a \pi_\theta(a|s) R(s, a) da ds$$

$$\nabla J = \int_s p(s) \int_a \nabla \pi_\theta(a|s) R(s, a) da ds$$

**Expectation is lost!**

**We don't know how to compute the gradient w.r.t. parameters**

# Training with Policy Gradient

**Problem:** we need gradients on parameters

$$J = \underset{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}}{E} R(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds$$

**Potential solution:** Finite differences

$$\nabla J \approx \frac{J_{\theta+\epsilon} - J_{\theta}}{\epsilon}$$

**Very noisy, especially if both J are sampled**

# Training with Policy Gradient

**Problem:** we need gradients on parameters

$$J = \underset{\substack{s \sim p(s) \\ a \sim \pi_{\theta}(s|a)}}{E} R(s, a) = \int_s p(s) \int_a \pi_{\theta}(a|s) R(s, a) da ds$$

**Wish list:**

- Analytical gradient
- Easy/stable approximations

Log-derivative trick

**Simple math question:**

$$\nabla \log \pi(z) = ? ? ?$$

Log-derivative trick

**Simple math question:**

$$\nabla \log \pi(z) = ? ? ?$$

$$\pi \cdot \nabla \log \pi(z) = \nabla \pi(z)$$

# Policy Gradient Approximation

*Algorithm: Monte-Carlo Policy Gradient (REINFORCE)*

$$\begin{aligned}\mathcal{J}(\theta) &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \pi(a|s; \theta) Q_{\pi}(s, a) \\ \nabla \mathcal{J}(\theta) &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \nabla \pi(a|s; \theta) Q_{\pi}(s, a) \\ &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \pi(a|s; \theta) \frac{\nabla \pi(a|s; \theta)}{\pi(a|s; \theta)} Q_{\pi}(s, a) && \text{by log-derivative trick} \\ &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \pi(a|s; \theta) \nabla \ln \pi(a|s; \theta) Q_{\pi}(s, a) \\ &= \mathbb{E}_{\pi_{\theta}}[\nabla \ln \pi(a|s; \theta) Q_{\pi}(s, a)] && \text{this expectation can be estimated by samples of episodes}\end{aligned}$$

# Policy Gradient

$$\nabla J = \int_s p(s) \int_a \nabla \pi_\theta(a|s) R(s, a) da ds$$

$$\nabla J \approx \frac{1}{N} \sum_{i=0}^N \nabla \log \pi_\theta(a|s) \cdot R(s, a)$$

# Supervised Learning vs. Policy Gradient

Supervised learning:

$$\nabla llh = E_{s, a_{opt} \sim D} \nabla \log \pi_{\theta}(a_{opt}|s)$$

Policy gradient:

$$\nabla J = E_{\substack{s \sim d(s) \\ a \sim \pi(a|obs(s))}} \nabla \log \pi_{\theta}(a|s) Q(s, a)$$

**Question: what is different? (apart from  $Q(s, a)$ )**



# Supervised Learning vs. Policy Gradient

Supervised learning:

$$\nabla llh = E_{s, a_{opt} \sim D} \nabla \log \pi_{\theta}(a_{opt}|s)$$

reference

Policy gradient:

$$\nabla J = E_{\substack{s \sim d(s) \\ a \sim \pi(a|obs(s))}} \nabla \log \pi_{\theta}(a|s) Q(s, a)$$

generated

# Supervised Learning vs. Policy Gradient

Supervised learning:

- Need (near-)optimal dataset
- Trains on reference sessions

Policy gradient:

- Need ~some data and reward function
- Trains on its own output

# Supervised Learning vs. Policy Gradient

## Supervised Learning

Need good reference ( $y_{opt}$ )

If model is *imperfect* [and **it is**],  
training:

$P(y_{next}|x, y_{prev\_ideal})$

prediction:

$P(y_{next}|x, y_{prev\_predicted})$

## Reinforcement Learning

Need reward function

Model learns to improve current policy. If policy is pure random, local improvements are unlikely to produce good translation.

# Supervised Learning vs. Policy Gradient

## Supervised Learning

- + Rather simple
- + Small variance
- Need good reference ( $y_{\text{opt}}$ )
- **Distribution shift:**  
different  $\mathbf{h}$  distribution  
when training vs generating

## Reinforcement learning

- + **Cold start problem**
- + Large variance (so far)
- Only needs  $x$  and  $r(s,a)$
- No **distribution shift**

# Supervised Learning vs. Policy Gradient

## Supervised Learning

pretraining

- + Rather simple
- + Small variance

- Need good reference ( $y_{opt}$ )
- **Distribution shift:**  
different  $h$  distribution  
when training vs generating

## Reinforcement learning

finetuning

- + **Cold start problem**
- + Large variance (so far)

- Only needs  $x$  and  $r(s,a)$
- No **distribution shift**

# Outro

This lecture covered:

- Exploration-vs-exploitation tradeoff in RL
- How to compare exploration strategies
- Algorithms:
  - Greedy, eps-greedy, softmax-sampling
  - Upper confidence bound based sampling
  - Probability matching and thompson sampling

# Outro

This lecture covered:

- Exploration-vs-exploitation tradeoff in RL
- How to compare exploration strategies
- Algorithms:
  - Greedy, eps-greedy, softmax-sampling
  - Upper confidence bound based sampling
  - Probability matching and thompson sampling
- We can solve **any structured prediction task** as RL problem
- Policy gradient -> optimize any reward (e.g. BLEU, Human Eval scores, CTR, ranking NDCG, etc.)

# Acknowledgements

This lecture uses materials from

- (1) [RL Lectures by David Siver](#) (licensed [CC-BY-NC 4.0](#))
- (2) [Practical RL lectures](#) by Yandex Data School ([Unlicense](#) license)



Questions?